

Introduction to Machine Learning 4:

Convolutional Neural Networks

All images were found under Creative Commons Licence or equivalent
(apart from the architecture images at the end)

Spot the Difference

(*Neamblysomus julianae*)



© Craig Jackson
Department of Biology
Norwegian University of Science and Technology (NTNU)
Trondheim, Norway



Male Bullock's Oriole; by Kevin Cole,
WikiMedia.Org, Creative Commons 2.0

(*Neamblysomus julianae*)



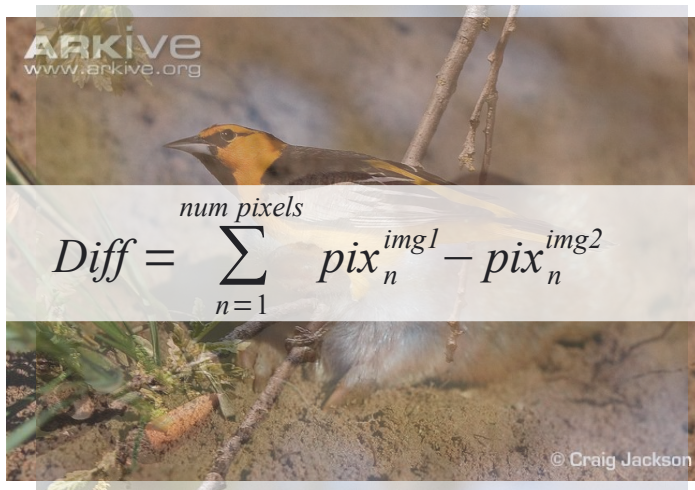
Take a
moment to
appreciate the
cuteness

© Craig Jackson
Department of Biology
Norwegian University of Science and Technology (NTNU)
Trondheim, Norway

Spot the Difference



Pixel by Pixel
comparison?

A photograph of a yellow and black bird, likely a species of warbler, perched on a thin branch. The bird has a bright yellow body with black markings on its wings and back. The background is a soft-focus blue sky. The ARKive logo and website address are visible in the top left corner, and the copyright notice '© Craig Jackson' is in the bottom right corner.
$$Diff = \sum_{n=1}^{num\ pixels} pix_n^{img1} - pix_n^{img2}$$

L1
distance

Spot the Difference



Pixel by Pixel
comparison?

$$Diff = \sum_{n=1}^{num\ pixels} pix_n^{img\ 1} - pix_n^{img\ 2}$$

L1
distance



Compare each image to all others,
calculate distances,

Then use a clustering algorithm to
perform classification: N-nearest
neighbours

Spot the Difference



Pixel by Pixel comparison / Nearest Neighbour

As you might have guessed this, doesn't really work that well...



Spot the Difference



Problems:



ARKive
www.arkive.org



© Craig Jackson









Image processing pixel weighting

We can perform image processing by looking at groups of pixels in an image and summing them up in different ways

For example we can cause a blur effect by averaging the values of neighboring pixels



Grey Value Pixels in Image

0.5	0.3	0.8	0.3
0.3	0.8	0.5	0.3
0.5	0.8	0.3	0.2
0.8	0.3	0.2	0.2

Grey Value Pixels in Image

0.5	0.3	0.8	0.3
0.3	0.8	0.5	0.3
0.5	0.8	0.3	0.2
0.8	0.3	0.2	0.2

We want to detect Patterns!

- Lines
- Edges
- Shapes
- Curves
- Etc...

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

“Kernel” or “Filter”

1	0	1
0	1	0
1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

“Kernel” or “Filter”

1	0	1
0	1	0
1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

The Kernel 'Slides' over the image multiplying the values of the image with the values in the Kernel. This is simple Matrix element-wise multiplication.

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

The Kernel 'Slides' over the image multiplying the values of the image with the values in the Kernel. This is simple Matrix element-wise multiplication.

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

The Kernel 'Slides' over the image multiplying the values of the image with the values in the Kernel. This is simple Matrix element-wise multiplication.

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1


The Kernel 'Slides' over the image multiplying the values of the image with the values in the Kernel. This is simple Matrix element-wise multiplication.

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Convolution Demo

Simplified Image




The 'Stride' is the amount of pixels the Kernel moves at each step, in this case, the Stride = 1

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image




The 'Stride' is the amount of pixels the Kernel moves at each step, in this case, the Stride = 1

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image




The 'Stride' is the amount of pixels the Kernel moves at each step, in this case, the Stride = 1

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Simplified Image



The 'Stride' is the amount of pixels the Kernel moves at each step, in this case, the Stride = 1

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Stride

- Can be changed depending on the size of the image or the patterns being looked for.

A large image might not need every pixel examined, since patterns might extend well beyond that scale. Kernel size can also play a role.

Simplified Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Kernel creates a new matrix, or convolved image

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1



“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

0 0x1	0 0x0	1 1x1	
0 1x0	0 0x1	0 1x0	
0 0x1	0 1x0	1 1x1	

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

Sum over the multiplication result			1
			1
			0
1	1	0	1

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

0 0x1	0 0x0	1 1x1	
0 1x0	0 0x1	0 1x0	
0 0x1	0 1x0	1 1x1	

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

0 0x1	0 0x0	1 1x1	
0 0x1	0 1x0	1 1x1	

SUM=2

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

0 <small>0x1</small>	0 <small>0x0</small>	1 <small>1x1</small>	
0 <small>1x0</small>	0 <small>1x1</small>	1 <small>1x1</small>	

SUM=2

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Convolution

2	

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

	0 0x1	0 0x0	1 1x1
	0 0x0	1 1x1	0 1x0
	1 1x1	0 0x0	0 0x1

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Convolution

2	

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

	0 0x1	0 0x0	1 1x1
	0 0x0	1 1x1	0 1x0
	1 1x1	0 0x0	0 0x1

SUM=3

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Convolution

2	3

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

1 1x1	0 0x0	1 1x1	
0 0x0	1 1x1	0 1x0	
1 1x1	0 1x0	0 0x1	

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Convolution

2	3

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

1 1x1	0 0x0	1 1x1	
0 0x0	1 1x1	0 1x0	
1 1x1	0 1x0	0 0x1	

SUM=4

"Kernel"
or
"Filter"

1	0	1
0	1	0
1	0	1

Convolution

2	3
4	

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

Multiplication Result

	0 _{0x1}	0 _{1x0}	1 _{1x1}
	0 _{1x0}	1 _{1x1}	0 _{0x0}
	1 _{1x1}	0 _{0x0}	1 _{1x1}

“Kernel”
or
“Filter”

1	0	1
0	1	0
1	0	1

Convolution

2	3
4	

Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

“Kernel”
or
“Filter”

Multiplication Result

	0 0x1	0 1x0	1 1x1
	1 1x0	0 0x0	0 0x1
	1 1x1	0 0x0	1 1x1

SUM=4


Convolution

2	3
4	4

Various Kernels Demo


Padding

- Using the method we have shown, convolutions will always be smaller than the input image,
- To change this, we can use Zero-Padding



Padding

- Using the method we have shown, convolutions will always be smaller than the input image,
- To change this, we can use Zero-Padding



Image

0	0	1	1
1	0	1	1
0	1	1	0
1	1	0	1

1	0	1
0	1	0
1	0	1

Convolution

2	3
4	4

Convolutions Stack...

Image

0	0	1	1	0	0	1	1
1	0	1	1	1	0	1	1
0	1	1	0	0	1	1	0
1	1	0	1	1	1	0	1
0	0	1	1	0	0	1	1
1	0	1	1	1	0	1	1
0	1	1	0	0	1	1	0
1	1	0	1	1	1	0	1

1	0	1
0	1	0
1	0	1

Convolution

2	3	2	3
4	4	4	4
2	3	2	3
4	4	4	4

1	0	1
0	1	0
1	0	1

Convolution

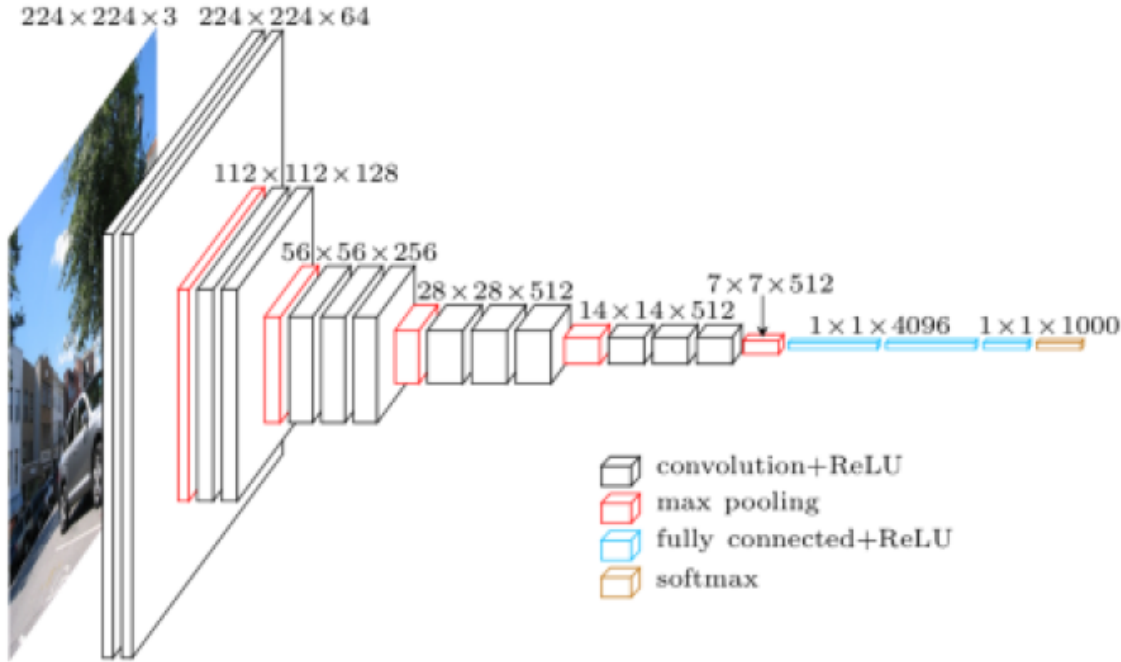
12	16
19	18

Kernels are 'Trained' along with the NN

- The values of the kernel matrix are adjusted along with the weights in the optimization process.

Visualizing Convolutional Networks filtered Images

Typical Architectures




VGG-16 model


Typically, a sequence of Convolutional layers followed by MaxPooling layers are used until a 1-dimensional layer is reached.

This is then used as an input to a some connected layers leading to a softmax classification output.


VGG-19



34-layer plain



34-layer residual



ResNet

Applications other than images

- Pattern recognition in temporal data
- 1D data
- Any data where there exist recurring patterns correlating to a class

Usefull Links

- <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>
- <https://cs231n.github.io/understanding-cnn/>
- https://github.com/keras-team/keras/blob/master/examples/conv_filter_visualization.py
- <https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html>